

Linux Fundamentals

Enrico RUSSO <enrico.russo@dibris.unige.it>

Preliminary remarks

- we take as a reference Linux **distribution** the Ubuntu 22.04 Long-term support (LTS), codename Jammy Jellyfish, **server version** (Ubuntu Server)
 - **distributions** are **customized versions** of Linux, and the most significant differences are
 - the **software package manager** (e.g., apt, dnf, pkg, ...)
 - **tools for configuring services** and keeping persistent their configurations (e.g., the network configuration manager)
 - some **paths** on the filesystem of the **tools/applications configuration files**
 - **server distros** are tailored for networks and services
 - they **do not include** a Graphical User Interface (**GUI**) and their installation involves a **minimal number of software packages**
 - they are deployed in datacenters (where servers run headless)

Preliminary remarks

- we administer each Linux server using the **shell** and a (virtual) **terminal**
 - the **shell** is a program that takes **commands** from the **keyboard** and gives them to the operating system to perform
 - many shell implementations exist but we take **bash** (or **sh**) as our reference
 - a **terminal** lets you **interact** with the **shell**
 - a **console** is generally a terminal in the **physical** sense, i.e., the primary terminal directly connected to a machine
 - a **virtual terminal** uses specific **network protocols** (e.g., telnet, ssh, ...) to connect to a remote machine and allow users to interact with its shell

man and tldr.sh

man command in Linux is used to display the user **manual** of any **command** that we can run on the terminal

- `man [command]` (e.g., `man sshd`)

tldr pages simplified and community-driven man pages.

- <https://tldr.sh/>

Users and accounts

Users and groups

Linux is a **multi-user** operating system and it can deal with several users simultaneously.

- each user needs an **account**, i.e., a login and (when required) a password
- users have a **personal environment** (e.g., a home directory, a shell, ...), which can be accessed only by them (and the system administrator and everyone knowing the password)
- inside the system the user is identified by the **user ID (UID)** and one or more **group IDs (GID)**

Accounts

There are three types of accounts on a Linux system:

- **root** account: is also called superuser and would have complete and unfettered control of the system (in recent distributions often does not have a password for interactive login)
- **system** accounts: are those needed for the operation of system-specific components
- **user** accounts: provide interactive access to the system for users and groups of users

User administration files

There are four main user administration files

- **/etc/passwd**: keeps the user account information
- **/etc/shadow**: holds the encrypted password of the corresponding account
- **/etc/group**: contains the group information for each account
- ~~**/etc/gshadow**: contains secure group account information~~

Becoming root

To become **root** user **from** an **unprivileged** (normal user account) account when no root password is set, we can use the *substitute user do* command

- **sudo -s** (enter the password of your Ubuntu user when requested)

According to its configuration file (/etc/sudoers), in Ubuntu distros users that are **members of the sudo group** can become root user.

```
# (%)user(/group) hostname=(runas-user:runas-group) command
# members of group sudo, logged in to any hostname, may run, as any user
# or group, any command
%sudo    ALL=(ALL:ALL) ALL
```

Manage accounts and groups

Command	Description
useradd	adds accounts to the system
usermod	modifies account attributes
passwd	changes user password
userdel	deletes accounts from system
groupadd	adds groups to the system
groupmod	modifies group attributes
groupdel	removes groups from the system
id	print user and group information of specified/current user

Filesystem

Filesystem Structure

Linux uses

- a **hierarchical** file system structure (much like an upside-down tree)
- with **root** (/) at the base of the file system and all other directories spreading from there
- **directories** have **specific purposes** and hold the same types of information following a hierarchy standard, namely **Filesystem Hierarchy Standard (FHS)**

File types

The following file-system objects can be found

- **normal** (text-)files
- **executable** files (binary files or shell scripts)
- **directories**
- **device** files: all physical devices (hard disks, DVD, USB, ...) are denoted by specific files
- symbolic or hard **links**: references to files
- **sockets**: used for inter-process communication (similar to TCP/IP sockets)
- (named) **pipes** (see later)

Filesystem Hierarchy Standard (FHS)

Path	Description
/etc	configuration files (disk configuration, valid user lists, groups, network configuration, hosts...)
/dev	device files: special files that provide an interface to a device driver (e.g., disk partitions, printers, and serial ports)
/bin	executable files available to all users
/sbin	executable files (usually for system administration)
/lib	shared library files
/usr	additional commands and data files

Filesystem Hierarchy Standard (FHS)

Path	Description
/var	variable-length files (e.g., log files).
/home	home directories
/boot	files for booting the system.
/tmp	temporary files.
/mnt	used to mount other temporary file systems.
/run	run-time variable data (e.g., running daemons)

FHS: /proc

/proc (process information pseudo-file system) is a **virtual** filesystem.

- it does not contain real files but **runtime system information** (e.g., system memory, devices mounted, hardware configuration, ...)
- a lot of system utilities are simply calls to files in this directory (e.g., `lsmod` prints `/proc/modules`)
- by altering files located in this directory (`/proc/sys` or `/sys`) you can even read/change kernel parameters (see also `sysctl` command) while the system is running (see `/etc/sysctl.conf`).

File information

While using `ls -lai` command, it displays various information related to file

inode	type	permissions	# links/dir	user	group	size	date	name
1984883	d	rw-rwxr-x	3	enrico	enrico	4096	set 25 13:59	.
1969658	d	rw-r-x---	24	enrico	enrico	4096	set 25 13:10	..
1984888	b	rw-rw-rw-	1	root	root	247,0	set 25 13:05	block_device
1984887	c	rw-rw-rw-	1	root	video	246,0	set 25 13:05	char_device
1984885	d	rw-rwxr-x	2	enrico	enrico	4096	set 25 13:53	dir
1984884	-	rw-rw-r--	2	enrico	enrico	10	set 25 13:53	file
1984884	-	rw-rw-r--	2	enrico	enrico	10	set 25 13:51	hard_link
1966437	-	rw-rw-r--	1	enrico	enrico	0	set 25 13:10	.hidden_file
1984886	l	rw-rwxrwx	1	enrico	enrico	4	set 25 13:03	link -> file
1984889	p	rw-rw-r--	1	enrico	users	0	set 25 13:06	pipe
1581	s	rw-rw-rw-	1	root	root	0	set 30 16:38	snapped.socket

File ownership/permissions

- Every file/directory **belongs** to a specific **user** or a **group** of users
- Every user/group may have **permissions** to **read**, **write**, and/or **execute**

user	group	others
rwx	rwx	rwx

	File	Directory
r	the file can be read	the directory's contents can be shown.
w	the file can be modified	the directory's contents can be modified (<u>requires the execute permission to be also set</u>)
x	the file can be executed	the directory can be accessed with cd

File ownership/permissions: commands

- `chown [user.group] [file]`: **changes ownership** of a file or a directory
- `chmod` **changes** the `rwX mode` bits of a file or directory
 - `+/ -`: **adds** or **removes** the mode bits
 - `u`: sets the permissions for the **owner**
 - `g`: sets the permissions for the **group** that of the owner belongs to
 - `o`: sets the permissions for the **other** users
 - `a`: sets the permissions for **all**

For example, `chmod g+w file` (add the write permission to the group owner of file)

	u	g	o							
type	permissions			# links/dir	user	group	size	date		name
d	rw	rw	r-x	2	enrico	enrico	4096	set 25	13:53	dir
-	rw	rw	r--	2	enrico	enrico	10	set 25	13:53	file

Special permissions

Run programs with **temporarily elevated privileges** in order to perform a specific task:

- user +s (or **SUID**): a file with SUID always **executes** as the **user who owns the file**, regardless of the user passing the command
- group +s (or **SGID**):
 - If set on a **file**, it allows the file to be **executed** as the **group that owns the file** (similar to SUID)
 - If set on a **directory**, any **files** created in the directory will have their **group ownership** set to that of the **directory owner**

Sticky bit:

- other +t: at the **directory level**, it **restricts file deletion**, i.e., Only the **owner** (and **root**) of a file can remove it within that directory

Filesystem commands and links

- create/remove file/dir: `touch`, `rm`, `rmdir`, `mkdir`, ...
- edit files: `nano`, `vi` (to exit: press `<Esc>`, Press `:` and use `q!` or `wq`)
- view files: `cat`, `less`, `more`, `head`, `tail` (use `tail -f` with logs), `grep`...
- find files: `find <options> <starting/path> <expression>`
 - `find / -name passwd`: find all files with name `passwd` starting from `/`
 - `find /etc -name passwd -exec wc -l {} \;`: find files with name `passwd` starting from `/etc` and for each found file count lines; use `{}` within the command to access the filename
 - `locate <filename>` (apt install `mlocate` and update the db with `updatedb`)
- `ln -s <path> <linkname>`: create a symbolic link

Mounting a file system

A file system must be **mounted** in order to be usable by the system.

- `mount`: see what is currently mounted (available for use) on the system.
- `mount -t <file_system_type> <device_to_mount> <directory_to_mount_to>`: mount a file system contained in the `<device_to_mount>` to the directory `<directory_to_mount_to>`
 - e.g., `mount -t iso9660 /dev/cdrom /mnt/cdrom` (mounting a cdrom).
- `umount <device/directory_mounted>`: unmount a filesystem
 - e.g., `umount /dev/cdrom`

Compress and extract files: tar (gzip/bzip)

- `tar -cvzf name-of-archive.tar.gz /path/to/directory-or-file1 .. /path/to/directory-or-fileN`: compress one (or multiple) entire directory or one (or multiple) file on Linux (c: create, v: display progress, z: compress with gzip, f: specify archive file name)
 - use j for compressing with bzip
 - `--exclude=/path/to/directory-or-file1`: exclude directory or file
- `tar -xvzf name-of-archive.tar.gz`: extract archive
 - `--strip 1`: strip off the first directory and extract the rest (`tar -xvz --strip 1 -f <archive>.tar.gz`)

Filesystem-related commands

Command	Description
<code>pwd</code>	prints current working directory
<code>ls</code>	lists the contents of a directory
<code>cd</code>	change the current path to the destination directory
<code>mkdir</code>	makes a new directory
<code>rmdir</code>	removes an empty directory
<code>cp</code>	copy file or directory
<code>mv</code>	move/rename file or directory

Filesystem-related commands

Command	Description
wc	word, line, character, and byte count
more	paging through text one screenful at a time
less	improved version of more allows backward/forward movement
head	display first lines of a file
tail	display last lines of a file
grep	print lines in a file matching a pattern

Processes

Processes

“On a UNIX system, everything is a file; if something is not a file, it is a process.”

— Machtelt Garrels, [Introduction To Linux: A Hands On Guide](#)

- Whenever you execute a program (a command), Linux starts a new **process**
- The operating system tracks processes through a unique five-digit ID number known as the **pid** (or the process ID)
- show processes:
 - `ps -f`: processes of current user
 - `ps -ef`: all processes, `ps -aux`: all processes, BSD style
 - `top`: realtime

Processes [2]

- **foreground** processes: by default, every process that you start runs in the foreground. It gets its **input** from the **keyboard** and sends its **output** to the **screen**
- **background** processes¹: a background process runs **without being connected to the terminal** (if it requires an input, it waits). Adding an ampersand & at the end of the command starts it as a background process
 - `jobs` (list background processes executed from the current shell)
 - `fg <jobid>`: put the job in foreground
 - `bg` or `CTRL-Z`: put the job in background
 - `kill %<jobid>` or `CTRL-C`: kill the job
- `kill <pid>`: kill a process (If a process ignores a regular `kill` command, add `-9`)

¹ **daemons**: processes that run in the background and are not interactive (they have no controlling terminal)

Background processes: example

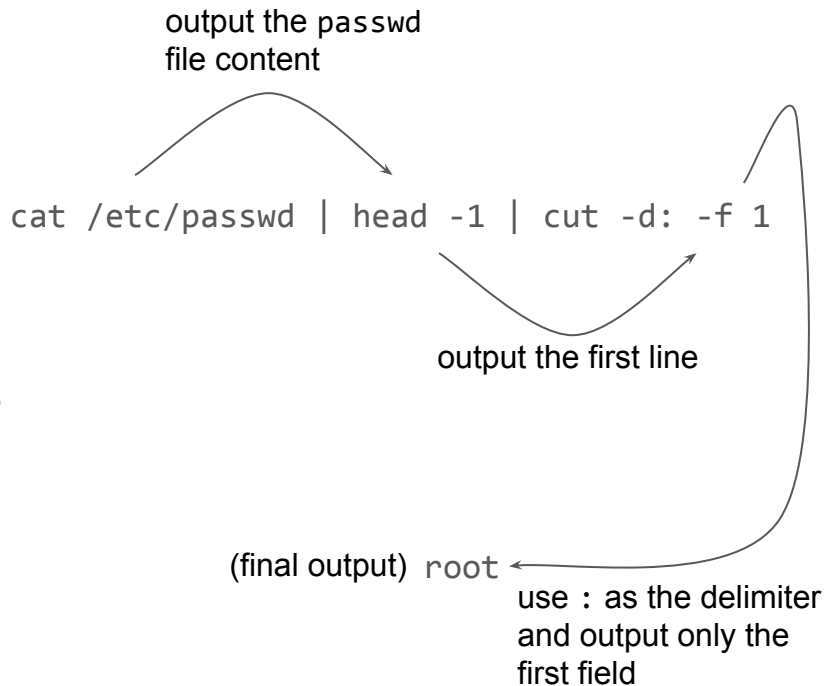
- `./loop.sh` (start in foreground and sends its output to the screen)
- CTRL-Z: put the job in background (but stopped)
[1]+ Stopped ./loop.sh
- `bg`: run the process in background (it sends its output to the screen)
[1]+ ./loop.sh &
- `jobs`: list jobs
[1]+ Running ./loop.sh &
- `kill %1`: kill the job [1]
[1]+ Terminated ./loop.sh

```
#!/bin/sh

# loop.sh
while true
do
    echo "hello"
    sleep 4
done
```

Pipes and Filters

- You can connect two commands together so that the output from one program becomes the input of the next program. Two or more commands connected in this way form a **pipe**
- To make a pipe, put a vertical bar `|` on the command line between two commands
- When a program takes its input from another program, it performs some operation on that input, and writes the result to the standard output. It is referred to as a **filter**
- **named pipes** provide such communication to processes using a special file (it is subjected to permission checks)



Packages and services

Advanced Package Tool (APT)

A packaging system is a way to provide programs and applications for installation (without building a program from source). Debian derivatives use the dpkg format and apt for interacting with the packaging system.

- `apt update`: updates the database of available packages
- `apt install <package-name>`: install <package-name>
- `apt search <str>`: search a package having <str> in the name or description
- `apt remove <package-name>`: remove <package-name>

Find the package providing a specific file: `dpkg -S <file>`, list packages: `dpkg -l`

systemd

systemd¹ is a Linux **initialization system** and **service manager**

- `systemctl`: show services status
- `systemctl start/stop/status <unit_name>`: start/stop/view a service

`journalctl` is a utility for **querying** and **displaying logs** from `journald`, the logging service of `systemd`.

- `journalctl -u <unit_name>`: message from specific unit (`journalctl --field _SYSTEMD_UNIT` list availables units)
- `journalctl path/to/executable`: message from specific executable
- `journalctl -f`: follow new messages (like `tail -f`)

¹<https://freedesktop.org/wiki/Software/systemd/>

Networking

Networking

`ip` command¹: show/manipulate routing, devices, policy routing and tunnels

- `ip a`: show addresses
- `ip r`: show routes
- `ip route add default via 192.168.53.2`: add default gw
- `ip addr add 192.168.53.100/24 dev ens33`: add new address
- `ip addr del 192.168.53.100/24 dev ens33`: delete address

Configurations with the `ip` command are not persistent!

¹https://access.redhat.com/sites/default/files/attachments/rh_ip_command_cheatsheet_1214_jcs_print.pdf

Netplan

netplan is a utility for easily configuring networking on a linux system (<https://netplan.io/>). It reads network configuration from `/etc/netplan/*.yaml` (see <https://netplan.io/examples/> for examples).

DHCP	Static
<pre>network: version: 2 ethernets: enp3s0: dhcp4: true</pre>	<pre>network: version: 2 ethernets: enp3s0: addresses: - 10.10.10.2/24 routes: - to: default via: 10.10.10.1 nameservers: search: [mydomain, otherdomain] addresses: [10.10.10.1, 1.1.1.1]</pre>

Test network configuration

- apply netconf configuration

- `netplan apply`

- check connectivity

- `ping [address]`

- check nameservers

- `host [name]`

Register names (locally)

The `/etc/hosts` is a plain text file that maps hostname to ip addresses.

```
127.0.0.1 localhost
```

```
127.0.1.1 vcc
```

```
192.168.58.2 gw gw.my.net
```

```
enrico@vcc:~$ ping gw.my.net
```

```
PING gw (192.168.58.2) 56(84) bytes of data.
```

```
64 bytes from gw (192.168.58.2): icmp_seq=1 ttl=128 time=0.179 ms
```

```
64 bytes from gw (192.168.58.2): icmp_seq=2 ttl=128 time=0.360 ms
```

```
64 bytes from gw (192.168.58.2): icmp_seq=3 ttl=128 time=0.397 ms
```

* Windows: `C:\Windows\System32\drivers\etc\hosts`

Remote terminal

SSH client

Windows

- openssh client
 - from **powershell** (as administrator): `Get-WindowsCapability -Online | ? Name -like 'OpenSSH*' (check) and Add-WindowsCapability -Online -Name OpenSSH.Client~~~~0.0.1.0 (install)`

connect with `ssh user@[address]`

- putty
 - <https://www.putty.org/>



SSH pubkey authentication

Public key authentication is a way of logging into a remote account using a cryptographic key rather than a password.

Client side (WINDOWS)

- generate a keypair: `ssh-keygen`
- copy the pub key to the server (using secure-copy from ssh): `scp id_rsa.pub enrico@192.168.58.100:.`

Server side (LINUX)

- create a `.ssh` directory (`mkdir ~/.ssh`)
- move `id_rsa.pub` in a new `~/.ssh/authorized_keys` file (`mv id_rsa.pub ~/.ssh/authorized_keys`)

Shell scripting

Shell scripting

A shell script is a program designed to be run by the Linux shell.

```
1.  #!/bin/bash
2.  # print current dir
3.  pwd
4.  # list files
5.  ls
6.  echo "end"
```

Line 1: **shebang** construct, specify the interpreter (sh shell)

Line 2,4: comments

Line 3,5: commands (listed in the order of execution)

Line 6: print the "end" string

Variables

The name of a variable can contain only letters (a to z or A to Z), numbers (0 to 9) or the underscore character (_).

1. `#!/bin/bash`
- 2.
3. `NAME="Enrico" # assign value`
4. `echo $NAME # print value (use $)`
5. `A=10`
6. `OP=$(expr $A '*' 2) # assign the output (man expr) to OP`
7. `TMP=`ls -l` # you can also use `` (backtick) instead of $()`

Environment variables

- Every shell has a set of attached variables
 - **system-defined** (e.g., \$PATH contains an ordered list of paths that Linux will search for executables when running a command)
 - **user-defined**: export command **promotes** a shell **variable** to an **environment variable**
- They are defined for the **d** and are inherited by any child shells or processes

Special variables

Name	Description
\$0	The filename of the current script
\$n	These variables correspond to the arguments with which a script was invoked. Here n is a positive decimal number corresponding to the position of an argument (the first argument is \$1, the second argument is \$2, and so on)
\$#	The number of arguments supplied to a script
\$*	All the arguments are double quoted. If a script receives two arguments, \$* is equivalent to "\$1 \$2"
\$@	All the arguments are <i>individually</i> double quoted. If a script receives two arguments, @\$ is equivalent to "\$1" "\$2"
\$?	The exit status of the last command executed
\$\$	The process number of the current shell. For shell scripts, this is the process ID under which they are executing
#!	The process number of the last background command

Operations

We can use external programs to perform basic (arithmetic, boolean, string, file test) operations.

- `expr` (man `expr`)
 - `expr 10 '/' 3`
- `test` (man `test`)
 - `test -f /etc/passwd; echo $?:` test if `/etc/passwd` is a file; the exit status is 0 for true
 - `test -d /tmp/testdir || mkdir /tmp/testdir:` create `/tmp/testdir` if it does not exist
 - `test -d /tmp/testdir && rmdir /tmp/testdir:` remove `/tmp/testdir` if it exists
 - `test 10 -gt 50; echo $?:` test if 10 is greater than 50

Redirect Input/Output

- `cat /etc/passwd > /tmp/users:` **redirects the output** of the command in the `/tmp/users` file
- `echo "test" >> /tmp/users:` **appends the output** in an existing file
- `wc -l < /tmp/users:` count the number of lines in the file by **redirecting the standard input** of the `wc` command from the file `/tmp/users`
- `command > /dev/null:` **discard the output**
- `command > /dev/null 2>&1:` **discard** both **output** of a command and its **error output** (2 represents `STDERR` and 1 represents `STDOUT`)
- `echo message 1>&2:` display a message on `STDERR` by **redirecting** `STDOUT` into `STDERR`

if..then..else

The base for the 'if' constructions is:

```
if [ expression ]; then  
    code if 'expression' is true.  
else  
    code if 'expression' is false.  
fi
```

```
1.  #!/bin/bash  
2.  a=100  
3.  if [ $a -gt 50 ]; then  
4.    echo "yes"  
5.  fi  
  
6.  b="enrico"  
7.  if [ $b == "Enrico" ]; then  
8.    echo "equal"  
9.  fi  
  
10. ls /tmp/nonexistent  
11. if [ $? == 0 ]; then  
12.   echo "yes"  
13. else  
14.   echo "no"  
15. fi
```

while

The base for the '**while**' constructions is:

```
while [ expression ];  
do  
  code if 'expression' is true.  
done
```

```
1.  #!/bin/bash  
2.  i=0  
3.  while [ $i -lt 10 ];  
4.  do  
5.    echo "i: $i"  
6.    i=$((expr $i '+' 1))  
7.  done
```

```
1.  #!/bin/bash  
  
2.  a=0  
3.  while true # infinite loop  
4.  do  
5.    echo "now: $a"  
6.    sleep 1  
7.    if [ -f /tmp/exit ]; then  
8.      echo "bye."  
9.      exit  
10.   fi  
11.   a=$((expr $a '+' 1))  
12. done
```

for

The base for the **'for'** constructions is:

```
for variable in [ expression ]  
do  
  code using $variable  
done
```

Some examples of [expression]:

- `1 2 3 4 5..N` (a numeric range)
- `string1 string2..stringN` (strings)
- `$(a_cmd_here)` (the output of a command)
- `{0..10..2}` (a range with a step)

```
1.  #!/bin/bash  
2.  for i in {0..10..2}  
3.    do  
4.      echo "Welcome $i times"  
5.    done  
  
6.  for e in $(ls -1 /etc)  
7.    do  
8.      if [ -d "/etc/$e" ]; then  
9.        echo "$e is a directory"  
10.     fi  
11.  done
```

for (example)

```
if [ $# -lt 1 ]; then
    echo "This command count the entries of a list of directories."
    echo
    echo "$0 [listofdir]"
    echo
    echo "Please specify a list of directories."
    exit
fi

for e in "$@"
do
    tmp=$(ls -1 $e | wc -l)
    echo "$e has $tmp entries."
done
```

Useful commands

sed	stream editor for filtering and transforming text
cut	cut out fields from `stdin` or files
tr	translate characters
sort	sort lines of text files
wc	count lines, words, or bytes
uniq	output the unique lines from the given input or (sorted) file
xargs	execute a command with piped arguments coming from another command
egrep	return lines that contain a pattern matching a given regular expression
logger	log a message to syslog

Exercises

1. create a script that accepts a filename *f*, a color *c* (red, blue, yellow, white or green) and an integer *i* as args and return *true* iff *c* appears *i* times in *f* (check if args are valid!).

A sample file follows:

```
red
blue
yellow
red
green
green
red
```

2. create a script that accepts an optional arg *a*. if *a* is the string 'empty' shows all the Linux accounts without password or show others otherwise.
3. create a script that accepts a program name *n* and a message *m* as args. It loops checking if *n* is running. If *n* is not running, logs for 3 times the message *m*.

Further Readings

- Ubuntu Server Guide (<https://ubuntu.com/server/docs>)
- The Linux Command Line (<http://linuxcommand.org/tlcl.php>)
- GNU Bash Manual (<https://www.gnu.org/software/bash/manual/>)
- Advanced Bash-Scripting Guide (<https://tldp.org/LDP/abs/html/>)
- 25 Free Books To Learn Linux For Free
(<https://itsfoss.com/learn-linux-for-free/>)